

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from plotly.offline import iplot, plot
from plotly.subplots import make_subplots
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import ConfusionMatrixDisplay
from sklearn.preprocessing import MinMaxScaler
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
import warnings
warnings.filterwarnings('ignore')

df = pd.read_csv("/content/train.csv")

df.sample(5)

print(f"Number of Row : {df.shape[0]}\nNumber of Columns : {df.shape[1]}")

```

➡ Number of Row : 2000
Number of Columns : 21

```

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   battery_power         2000 non-null   int64
 1   blue                  2000 non-null   int64
 2   clock_speed           2000 non-null   float64
 3   dual_sim              2000 non-null   int64
 4   fc                    2000 non-null   int64
 5   four_g                2000 non-null   int64
 6   int_memory            2000 non-null   int64
 7   m_dep                 2000 non-null   float64
 8   mobile_wt             2000 non-null   int64
 9   n_cores               2000 non-null   int64
10   pc                    2000 non-null   int64
11   px_height             2000 non-null   int64
12   px_width              2000 non-null   int64
13   ram                   2000 non-null   int64
14   sc_h                  2000 non-null   int64
15   sc_w                  2000 non-null   int64
16   talk_time             2000 non-null   int64
17   three_g              2000 non-null   int64
18   touch_screen          2000 non-null   int64
19   wifi                  2000 non-null   int64
20   price_range           2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB

```

```

df.isna().sum()

battery_power    0
blue             0
clock_speed      0
dual_sim         0
fc               0
four_g           0
int_memory       0
m_dep            0
mobile_wt        0
n_cores          0
pc               0
px_height        0
px_width         0
ram              0
sc_h             0
sc_w             0
talk_time        0
three_g          0
touch_screen     0
wifi             0
price_range      0
dtype: int64

```

```
df.describe()
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores
count	2000.000000	2000.0000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000	2000.000000
mean	1238.518500	0.4950	1.522250	0.509500	4.309500	0.521500	32.046500	0.501750	140.249000	4.520500
std	439.418206	0.5001	0.816004	0.500035	4.341444	0.499662	18.145715	0.288416	35.399655	2.287837
min	501.000000	0.0000	0.500000	0.000000	0.000000	0.000000	2.000000	0.100000	80.000000	1.000000
25%	851.750000	0.0000	0.700000	0.000000	1.000000	0.000000	16.000000	0.200000	109.000000	3.000000
50%	1226.000000	0.0000	1.500000	1.000000	3.000000	1.000000	32.000000	0.500000	141.000000	4.000000
75%	1615.250000	1.0000	2.200000	1.000000	7.000000	1.000000	48.000000	0.800000	170.000000	7.000000
max	1998.000000	1.0000	3.000000	1.000000	19.000000	1.000000	64.000000	1.000000	200.000000	8.000000

8 rows × 11 columns


```
df.duplicated().any()
```

False

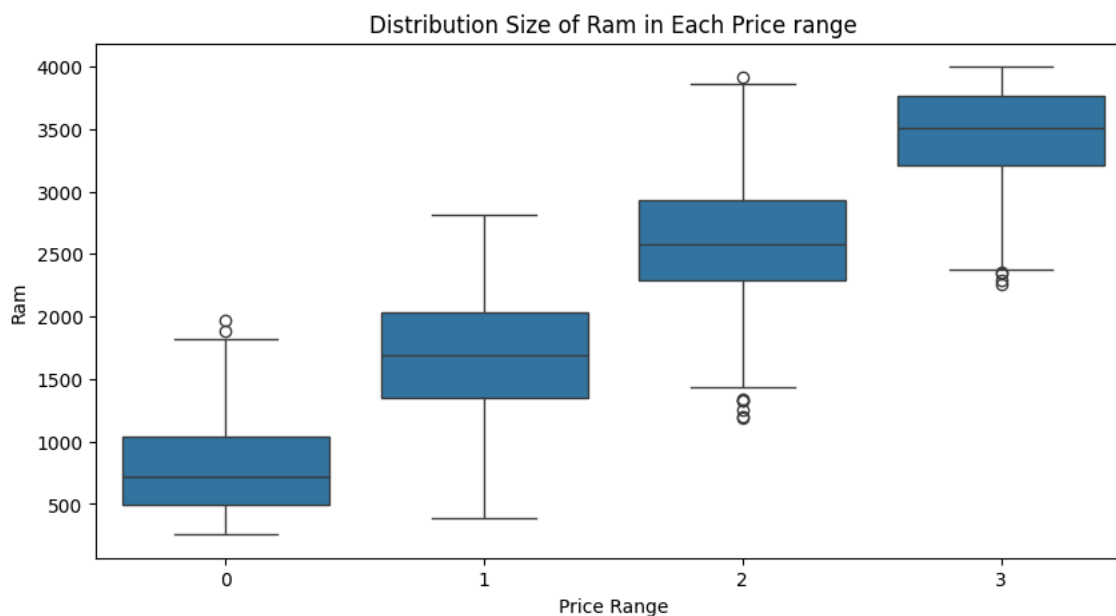
```
df_battery_price = df.groupby('price_range')['battery_power'].mean()
```

```
ipplot(px.line(df_battery_price,
               labels={'value':'Mean of Battery Capacity','price_range':'Price Range'},
               color_discrete_sequence=['red']
               ))
```

```
plt.figure(figsize=(10,5))
plt.title('Distribution Battery Power in Each Price range')
sns.boxplot(x=df['price_range'],y=df['battery_power'])
plt.xlabel('Price Range')
plt.ylabel('Battery Power')
plt.show()
```



```
plt.figure(figsize=(10,5))
plt.title('Distribution Size of Ram in Each Price range')
sns.boxplot(x=df['price_range'],y=df['ram'])
plt.xlabel('Price Range')
plt.ylabel('Ram')
plt.show()
```



```
df_4g = df['four_g'].value_counts()
iplot(px.pie(values=df_4g,
              names=['Support 4G', 'Not Support 4G'],
              template='plotly_dark',
              title='Is Support 4G ?'
              ).update_traces(textinfo='label+percent'))

df_3g = df['three_g'].value_counts()
iplot(px.pie(values=df_3g,
              names=['Support 3G', 'Not Support 3G'],
              template='plotly_dark',
              title='Is Support 3G ?'
              ).update_traces(textinfo='label+percent'))

x = df.drop(columns='price_range')
y = df.price_range

scaler = MinMaxScaler()
x = scaler.fit_transform(x)

x_train , x_test , y_train , y_test = train_test_split(x,y,test_size=0.2)

print(f'Shape of X_Train {x_train.shape}')
print(f'Shape of X_Test {x_test.shape}')
print(f'Shape of Y_Train {y_train.shape}')
print(f'Shape of Y_Test {y_test.shape}')

Shape of X_Train (1600, 20)
Shape of X_Test (400, 20)
Shape of Y_Train (1600,)
Shape of Y_Test (400,)
```

```
model_params = {
    'svm':{
        'model' : SVC(gamma='auto'),
        'params':{
            'C':[1,10,20],
            'kernel':['rbf','linear']
        }
    },
    'random_forest':{
        'model':RandomForestClassifier(),
        'params':{
            'n_estimators':[1,5,10]
        }
    },
    'logistic_regression':{
        'model':LogisticRegression(solver='liblinear',multi_class='auto'),
        'params':{
            'C':[1,5,10]
        }
    }
}
```

```
scores = []

for model_name , mp in model_params.items():
    clf = GridSearchCV(mp['model'],mp['params'],cv=5,return_train_score=False)
    clf.fit(x,y)
    scores.append({
        'model':model_name,
        'best_scores':clf.best_score_,
        'best_params':clf.best_params_
    })

pd.DataFrame(scores,columns=['model','best_scores','best_params'])
```

	model	best_scores	best_params
0	svm	0.9675	{'C': 20, 'kernel': 'linear'}
1	random_forest	0.8080	{'n_estimators': 10}
2	logistic_regression	0.8375	{'C': 10}

```
model_svm = SVC(kernel='linear',C=20)
model_svm.fit(x_train,y_train)
```

▼

SVC

SVC(C=20, kernel='linear')


```
score_svm_train = model_svm.score(x_train,y_train)
print(f"Train accuracy: {score_svm_train}")
```

Train accuracy: 0.976875

```
score_svm_test = model_svm.score(x_test,y_test)
print(f"Test accuracy: {score_svm_test}")
```

Test accuracy: 0.9625

```
ConfusionMatrixDisplay.from_estimator(model_svm,
                                      x_test,
                                      y_test);
```



```
model_LR = LogisticRegression(C=10)
model_LR.fit(x_train,y_train)
```

```
▼ LogisticRegression
LogisticRegression(C=10)
```


```
score_LR_train = model_LR.score(x_train,y_train)
print(f"Train accuracy: {score_LR_train}")
```

Train accuracy: 0.974375

```
score_LR_test = model_LR.score(x_test,y_test)
print(f"Test accuracy: {score_LR_test}")
```

Test accuracy: 0.965

```
ConfusionMatrixDisplay.from_estimator(model_LR,
                                      x_test,
                                      y_test);
```



```
model_RFC = RandomForestClassifier(n_estimators=10,random_state=42)
model_RFC.fit(x_train,y_train)
```

```
▼ RandomForestClassifier
RandomForestClassifier(n_estimators=10, random_state=42)
```

```
score_RFC_train = model_RFC.score(x_train,y_train)
print(f"Train accuracy: {score_RFC_train}")
```

Train accuracy: 0.995625

```
score_RFC_test = model_RFC.score(x_test,y_test)
print(f"Test accuracy: {score_RFC_test}")
```

Test accuracy: 0.7625

```
from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
import xgboost as xgb
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
# Training Gradient Boosting model
gb_model = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
gb_model.fit(x_train, y_train)
# Training XGBoost model
xgb_model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)
xgb_model.fit(x_train, y_train)
# Making predictions on the testing set
gb_predictions = gb_model.predict(x_test)
xgb_predictions = xgb_model.predict(x_test)
```

```
mae = mean_absolute_error(y_test, gb_predictions)
print("Mean Absolute Error (MAE):", mae)
# Calculate Mean Squared Error (MSE)
mse = mean_squared_error(y_test, gb_predictions)
print("Mean Squared Error (MSE):", mse)
# Calculate Root Mean Squared Error (RMSE)
rmse = mean_squared_error(y_test, gb_predictions, squared=False)
print("Root Mean Squared Error (RMSE):", rmse)
# Calculate R-squared (R2) score
r2 = r2_score(y_test, gb_predictions)
print("R-squared (R2) Score:", r2)
```

Mean Absolute Error (MAE): 0.22347228527141083
Mean Squared Error (MSE): 0.07961726983879017